

Lecture #2

Frequent Pattern Mining: Itemsets, Association Rules, Algorithms, and Applications*

* Slides partly based on Jiawei Han, and Jeff Ullman

Frequent Pattern Mining: Overview

- **Frequent pattern**

- ▲ a pattern (a set of items, subsequences, substructures, etc.) that occurs frequently in a data set

- **Motivation:** Finding inherent regularities in data

- ▲ What products were often purchased together in supermarket?

- ▲ What are the subsequent purchases after buying a PC?

- ▲ What kinds of DNA are sensitive to this new drug?

Frequent Pattern Mining: Overview

- **Applications**

- ▶ Market-Basket analysis to improve sales
- ▶ Web log (click stream) analysis
- ▶ DNA sequence analysis
- ▶ Plagiarism detection
- ▶ ...

- First proposed by Agrawal, Imielinski, and Swami [1993] in the context of frequent itemsets and association rule mining

The Market-Basket Model

- A large set of *items*, e.g., things sold in a supermarket
- A large set of *baskets*, each of which is a “small” set of the items, e.g., the things one customer buys on one day
 - ▲ baskets are also referred to as *transactions* in the literature

Frequent Itemsets: Concept

- **Simple computational question:** find sets of items that appear “frequently” in the baskets.
- **Support** for itemset I
 - ▲ Absolute: the number of baskets containing all items in I
 - ▲ Relative: the fraction of baskets that contain items in I (i.e., the probability that a basket contains items in I)
- **Frequent Itemset**
 - ▲ Given a support threshold s , a set of items appearing in at least s baskets is called a *frequent itemset*

Frequent Itemsets: Example

- Items={milk, coke, pepsi, beer, juice}.
- Support = 3 baskets.
- Example baskets with items.

$$B_1 = \{m, c, b\}$$

$$B_2 = \{m, p, j\}$$

$$B_3 = \{m, b\}$$

$$B_4 = \{c, j\}$$

$$B_5 = \{m, p, b\}$$

$$B_6 = \{m, c, b, j\}$$

$$B_7 = \{c, b, j\}$$

$$B_8 = \{b, c\}$$

- Frequent itemsets: {m}, {c}, {b}, {j},
{m,b} , {b,c} , {c,j}

Frequent Itemsets: Applications

- “Classic” application was analyzing what people bought together in a brick-and-mortar store
 - ▲ Apocryphal story of “diapers and beer” discovery
 - ▲ Used to position potato chips between diapers and beer to enhance sales of potato chips.
- Many other applications, including plagiarism detection
 - ▲ Items = documents; baskets = sentences
 - ▲ Basket/sentence contains all the items/documents that have that sentence
 - ▲ MOSS software: <https://theory.stanford.edu/~aiken/moss/>

Frequent Patterns: Combinatorial Explosion and ``Closed'' Patterns

- A long pattern contains a combinatorial number of sub-patterns, e.g., $\{a_1, \dots, a_{100}\}$ contains $\binom{100}{1} + \binom{100}{2} + \dots + \binom{100}{100} = 2^{100} - 1 = 1.27 \times 10^{30}$ sub-patterns!
- Solution: Mine ``closed'' patterns
- An itemset I is closed if I is *frequent* and there exists *no super-set* $J \supset I$, with the same support as I
- Closed pattern is a lossless compression of frequent patterns: Reducing the # of patterns and rules

Association Rules: Concept

- If-then rules about the contents of baskets
- $\{i_1, i_2, \dots, i_k\} \rightarrow j$ means: “if a basket contains all of i_1, \dots, i_k then it is *likely* to contain j ”
 - ▲ Example: {bread, peanut-butter} \rightarrow jelly
- *Confidence* of this association rule is the “probability” of j given i_1, \dots, i_k
 - ▲ That is, the fraction of the baskets with i_1, \dots, i_k that also contain j

Subtle point: “probability” implies there is a process generating random baskets. Really we’re just computing the fraction of baskets, because we’re computer scientists, not statisticians.

Association Rules: Example

- + $B_1 = \{m, c, b\}$
- $B_3 = \{m, b\}$
- $B_5 = \{m, p, b\}$
- $B_7 = \{c, b, j\}$
- $B_2 = \{m, p, j\}$
- $B_4 = \{c, j\}$
- + $B_6 = \{m, c, b, j\}$
- $B_8 = \{b, c\}$

- An association rule: $\{m, b\} \rightarrow c$
 - ▲ Confidence = $2/4 = 50\%$

Computation Model

- Typically, data is a file consisting of a list of baskets
- The true cost of mining disk-resident data is usually the **number of disk I/O's**
- In practice, we read the data in *passes* – all baskets read in turn.
 - ▲ Thus, we measure the cost by the **number of passes** an algorithm takes

Main-Memory Bottleneck

- For many frequent-itemset mining algorithms, main memory is the critical resource
 - As we read baskets, we need to count something, e.g., occurrences of pairs of items
- The number of different things we can count is limited by main memory
 - ▲ Swapping counts in/out is a disaster

Finding Frequent Pairs


- The hardest problem often turns out to be **finding the frequent pairs**
 - ▲ Why? Often frequent pairs are common, frequent triples are rare
 - ▲ Support threshold is usually set high enough that you don't get too many frequent itemsets
- We'll first concentrate on computing frequent pairs, and then extend to larger sets.

Finding Frequent Pairs: Naïve Algorithm

- Read file once, counting in main memory the occurrences of each pair
 - ▶ From each basket of n items, generate its $n(n-1)/2$ pairs by two nested loops
- Fails if $(\text{\#items})^2$ exceeds main memory
 - ▶ Example 1: Walmart sells 100K items, so probably OK.
 - ▶ Example 2: Web has 100B pages, so definitely not OK.

Counting in Main Memory: Two Approaches


- Approach 1: Count all pairs using a **triangular matrix**
 - ▲ `Count[i,j]` in row `i`, column `j`, provided $i < j$
 - ▲ use a “ragged array,” so the empty triangle is not there (more details later)



Note: assume integers are 4 bytes

Counting in Main Memory: Two Approaches

- Approach 2: Store a table of pairs with count > 0
 - ▲ Keep a table of triples $[i, j, c]$ = “the count of the pair of items $\{i, j\}$ is c ”



Tabular method

Note: assume integers are 4 bytes

Triangular Matrix vs. Tabular Method

- **Triangular matrix**
 - ▲ Requires 4 bytes per pair
 - **Tabular method**
 - ▲ Requires 12 bytes per pair, but only for those pairs with count > 0
- **Comparison:** tabular approach beats triangular matrix only when **at most 1/3 of all pairs have a nonzero count**

Triangular Matrix as One-Dimensional Array

- Number items $1, 2, \dots, n$
 - ▲ Requires table of size $O(n)$ to convert item names to consecutive integers
- Count $\{i, j\}$ only if $i < j$
 - ▲ Keep pairs in the order $\{1,2\}, \{1,3\}, \dots, \{1,n\}, \{2,3\}, \{2,4\}, \dots, \{2,n\}, \{3,4\}, \dots, \{3,n\}, \dots, \{n-1,n\}$
- Find pair $\{i, j\}$, where $i < j$, at the position:

$$(i - 1)(n - i/2) + j - i$$
- Total number of pairs $n(n-1)/2$
 - ▲ total bytes about $2n^2$

The A-Priori Algorithm

- Monotonicity of “Frequent” Candidate Pairs
- Extension to Larger Itemsets

A-Priori Algorithm

- A two-pass approach called “*a-priori*” limits the need for main memory

- Key idea: *monotonicity*: if a set of items appears at least s times, so does every subset of the set

- *Contrapositive for pairs*: if item i does not appear in s baskets, then no pair including i can appear in s baskets

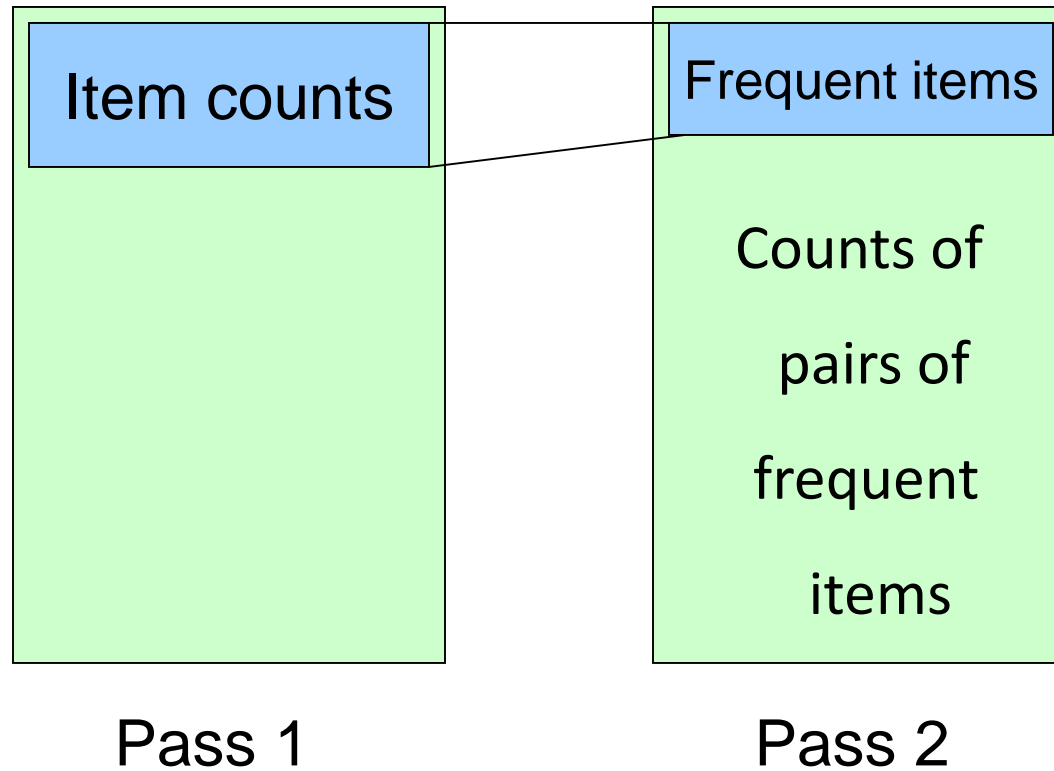
A-Priori Algorithm (2)

- **Pass 1:** Read baskets and count in main memory the occurrences of each item.
 - ▲ Requires only memory proportional to #items
- Items that appear at least s times are the *frequent items*

A-Priori Algorithm (3)

- **Pass 2**: Read baskets again and count in main memory only those pairs both of which were found in Pass 1 to be frequent
- Requires memory proportional to square of *frequent items* only (for counts), plus a table of the frequent items (so you know what must be counted).

Picture of A-Priori



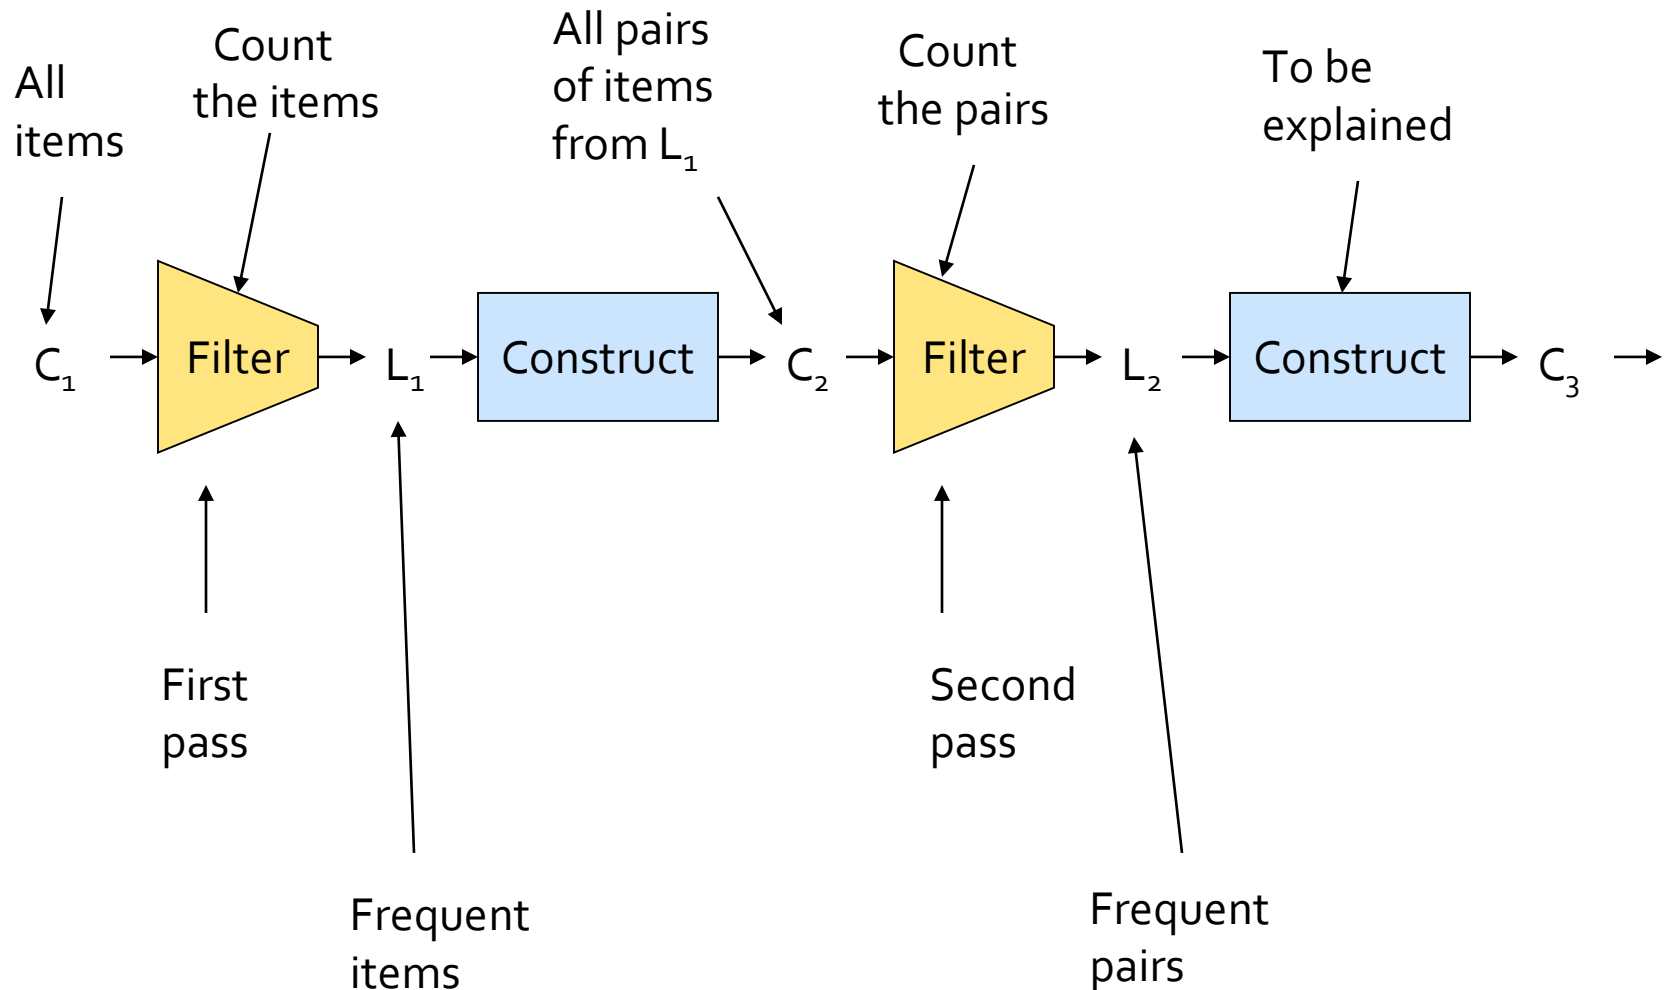
Detail for A-Priori

- You can use the triangular matrix method with n = number of frequent items
 - ▲ May save space compared with storing triples
- **Trick:** number frequent items 1, 2,... and keep a table relating new numbers to original item numbers

Frequent Triples etc.

- For each size of itemsets k , we construct two sets of k -sets (sets of size k):
 - ▲ C_k = *candidate k -sets* = those that might be frequent sets (support $\geq s$) based on information from the pass for itemsets of size $k - 1$
 - ▲ L_k = the set of *truly frequent k -sets*

A-Priori Algorithm: Pictorial Illustration



Passes Beyond Two

- C_1 = all items
- In general, L_k = members of C_k with support $\geq s$
 - ▶ Requires one pass
- C_{k+1} = $(k+1)$ -sets, each k of which is in L_k
- How would you generate C_{k+1} from L_k ?
 - ▶ Enumerating all sets of size $k+1$ and testing each seems really dumb

A-Priori Algorithm: Example

$\text{Sup}_{\min} = 2$

Database TDB

Bid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E

1st scan

C_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

L_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{E}	3

L_2

Itemset	sup
{A, C}	2
{B, C}	2
{B, E}	3
{C, E}	2

C_2

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

2nd scan

C_2

Itemset
{A, B}
{A, C}
{A, E}
{B, C}
{B, E}
{C, E}

C_3

Itemset
{B, C, E}

3rd scan

L_3

Itemset	sup
{B, C, E}	2

A-Priori Algorithm: Pseudocode

C_k : Candidate itemset of size k

L_k : frequent itemset of size k

$L_1 = \{\text{frequent items}\};$

for ($k = 1; L_k \neq \emptyset; k++$) **do begin**

C_{k+1} = candidates generated from L_k ;

for each basket t in database do

increment the count of all candidates in C_{k+1} that
are contained in t

L_{k+1} = candidates in C_{k+1} with min_support

end

return $\cup_k L_k$;

Memory Requirements

- At the k^{th} pass, you need space to count each member of C_k
- In realistic cases, because you need fairly high support, the number of candidates of each size drops, once you get beyond pairs

Improvements over A-Priori Algorithm

- **The PCY (Park-Chen-Yu) Algorithm**

- ▶ Improvement to A-Priori

- ▶ Exploits empty memory on first Pass

- ▶ Frequent buckets

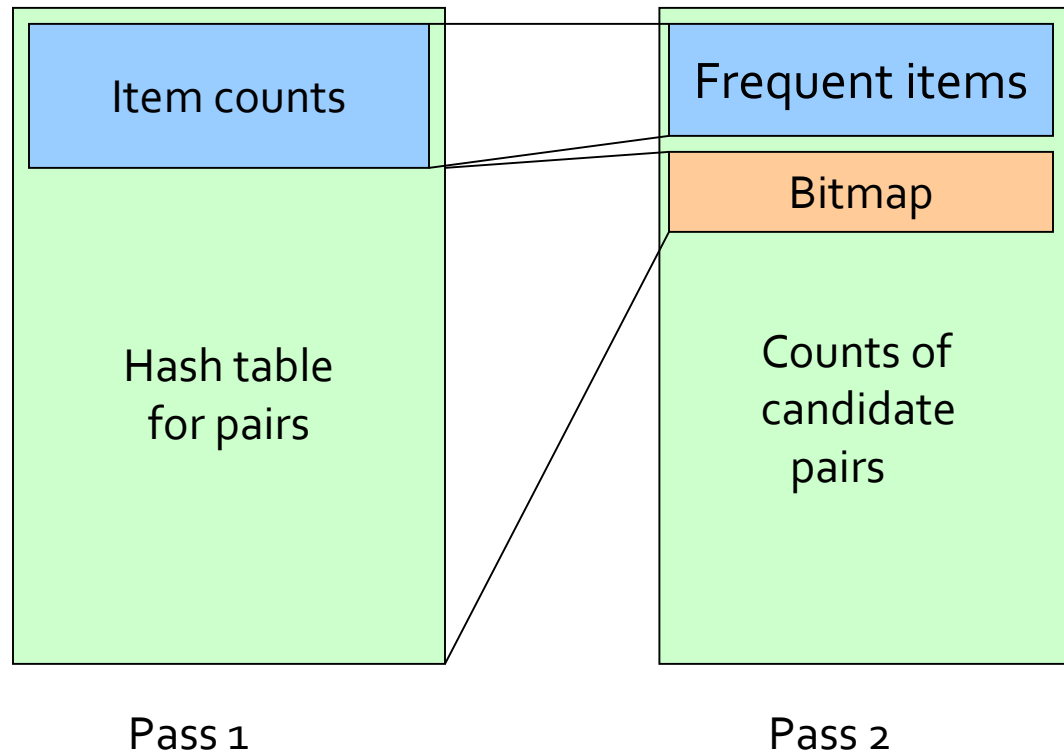
PCY Algorithm

- During Pass 1 of A-priori, most memory is idle
- Use that memory to keep counts of buckets into which pairs of items are hashed
 - ▲ Just the count, not the pairs themselves
- For each basket, enumerate all its pairs, hash them, and increment the resulting bucket count by 1

PCY Algorithm (2)

- A bucket is *frequent* if its count is at least the support threshold
- If a bucket is not frequent, no pair that hashes to that bucket could possibly be a frequent pair
- On Pass 2, we only count pairs of frequent items that also hash to a frequent bucket.
- A *bitmap* tells which buckets are frequent, using only one bit per bucket (i.e., $1/32$ of the space used on Pass 1).

Picture of PCY



Pass 1 of PCY: Memory Organization

- Space to count each item
 - ▲ One (typically) 4-byte integer per item
- Use the rest of the space for as many integers, representing buckets, as we can

PCY Algorithm: Pass 1

```
FOR (each basket) {  
    FOR (each item in the basket)  
        add 1 to item's count;  
    FOR (each pair of items) {  
        hash the pair to a bucket;  
        add 1 to the count for that bucket  
    }  
}
```

Observations about Buckets

1. A bucket that a frequent pair hashes to is surely frequent
 - ▲ We cannot eliminate any member of this bucket
2. Even without any frequent pair, a bucket can be frequent
 - ▲ Again, nothing in the bucket can be eliminated
3. But if the count for a bucket is less than the support s , all pairs that hash to this bucket can be eliminated, even if the pair consists of two frequent items

PCY Algorithm: Between Passes

- Replace the buckets by a bit-vector (the “**bitmap**”):
 - ▲ 1 means the bucket is frequent; 0 means it is not
- Also, decide which items are frequent and list them for the second pass

PCY Algorithm: Pass 2

- Count all pairs $\{i, j\}$ that meet the conditions for being a **candidate pair**:
 1. Both i and j are frequent items
 2. The pair $\{i, j\}$, hashes to a bucket number whose bit in the bit vector is 1

PCY Algorithm: Memory Details

- Buckets require a few bytes each
 - ▲ Note: we don't have to count past s
 - If $s < 2^{16}$, 2 bytes/bucket will do
 - ▲ # buckets is $O(\text{main-memory size})$
- On second pass, a table of (item, item, count) triples is essential
 - ▲ Thus, hash table on Pass 1 must eliminate 2/3 of the candidate pairs for PCY to beat a-priori

All (Or Most) Frequent Itemsets In ≤ 2 Passes

- Simple Algorithm
- Savasere-Omiecinski- Navathe (SON) Algorithm
- Toivonen's Algorithm

Simple Algorithm

- Take a random sample of the market baskets
 - ▲ **Do not sneer**; “random sample” is often a cure for the problem of having too large a dataset
- Run a-priori or one of its improvements (for sets of all sizes, not just pairs) in main memory, so you don't pay for disk I/O each time you increase the size of itemsets.
- Use as your support threshold a suitable, scaled-back number
 - ▲ **Example**: if your sample is $1/100$ of the baskets, use $s/100$ as your support threshold instead of s .

Simple Algorithm: Option

- Optionally, verify that your guesses are truly frequent in the entire data set by a second pass
- But you don't catch sets frequent in the whole but not in the sample
 - ▲ Smaller threshold, e.g., $s/125$ instead of $s/100$, helps catch more truly frequent itemsets (requires more space)

SON Algorithm: Pass 1

- Partition the baskets into small subsets
- Read each subset into main memory and perform the first pass of the simple algorithm on each subset
 - ▶ Parallel processing of the subsets a good option
- An itemset is a candidate if it is frequent (with support threshold suitably scaled down) in *at least one* subset.

SON Algorithm: Pass 2

- On a second pass, count all the candidate itemsets and determine which are frequent in the entire set
- **Key “monotonicity” idea**: an itemset cannot be frequent in the entire set of baskets unless it is frequent in at least one subset

Toivonen's Algorithm

- Start as in the simple algorithm, but lower the threshold slightly for the sample
 - ▲ Example: if the sample is 1% of the baskets, use $s/125$ as the support threshold rather than $s/100$
 - ▲ Goal is to avoid missing any itemset that is frequent in the full set of baskets


Toivonen's Algorithm (2)

- Add to the itemsets that are frequent in the sample the *negative border* of these itemsets
- An itemset is in the negative border if it is not deemed frequent in the sample, but all its immediate subsets are
 - ▲ *Immediate subset* = “delete exactly one element”

Example: Negative Border

- $\{A,B,C,D\}$ is in the negative border if and only if:
 1. It is not frequent in the sample, but
 2. All of $\{A,B,C\}$, $\{B,C,D\}$, $\{A,C,D\}$, and $\{A,B,D\}$ are
- $\{A\}$ is in the negative border if and only if it is not frequent in the sample
 - ▲ Because the empty set is always frequent. Unless there are fewer baskets than the support threshold (silly case)
- Useful trick: When processing the sample by A-Priori, each member of C_k is either in L_k or in the negative border, never both

Picture of Negative Border




Toivonen's Algorithm (3)

- In a second pass, count all candidate frequent itemsets from the first pass, and also count sets in their negative border
- If no itemset from the negative border turns out to be frequent, then the candidates found to be frequent in the whole data are *exactly* the frequent itemsets

Toivonen's Algorithm (4)

- What if we find that something in the negative border is actually frequent?
- We must start over again with another sample!
- Try to choose the support threshold so the probability of failure is low, while the number of itemsets checked on the second pass fits in main-memory

If Something in the Negative Border is Frequent ...



Theorem

- If there is an itemset that is frequent in the whole, but not frequent in the sample, then there is a member of the negative border for the sample that is frequent in the whole

Proof

- Suppose not; i.e.;
 1. There is an itemset S frequent in the whole but not frequent in the sample, and
 2. Nothing in the negative border is frequent in the whole
- Let T be a **smallest** subset of S that is not frequent in the sample
- T is frequent in the whole (S is frequent + monotonicity)
- T is in the negative border (else not “smallest”)