

Lecture #3: Machine Learning

Janardhan Rao (Jana) Doppa

School of EECS, Washington State University

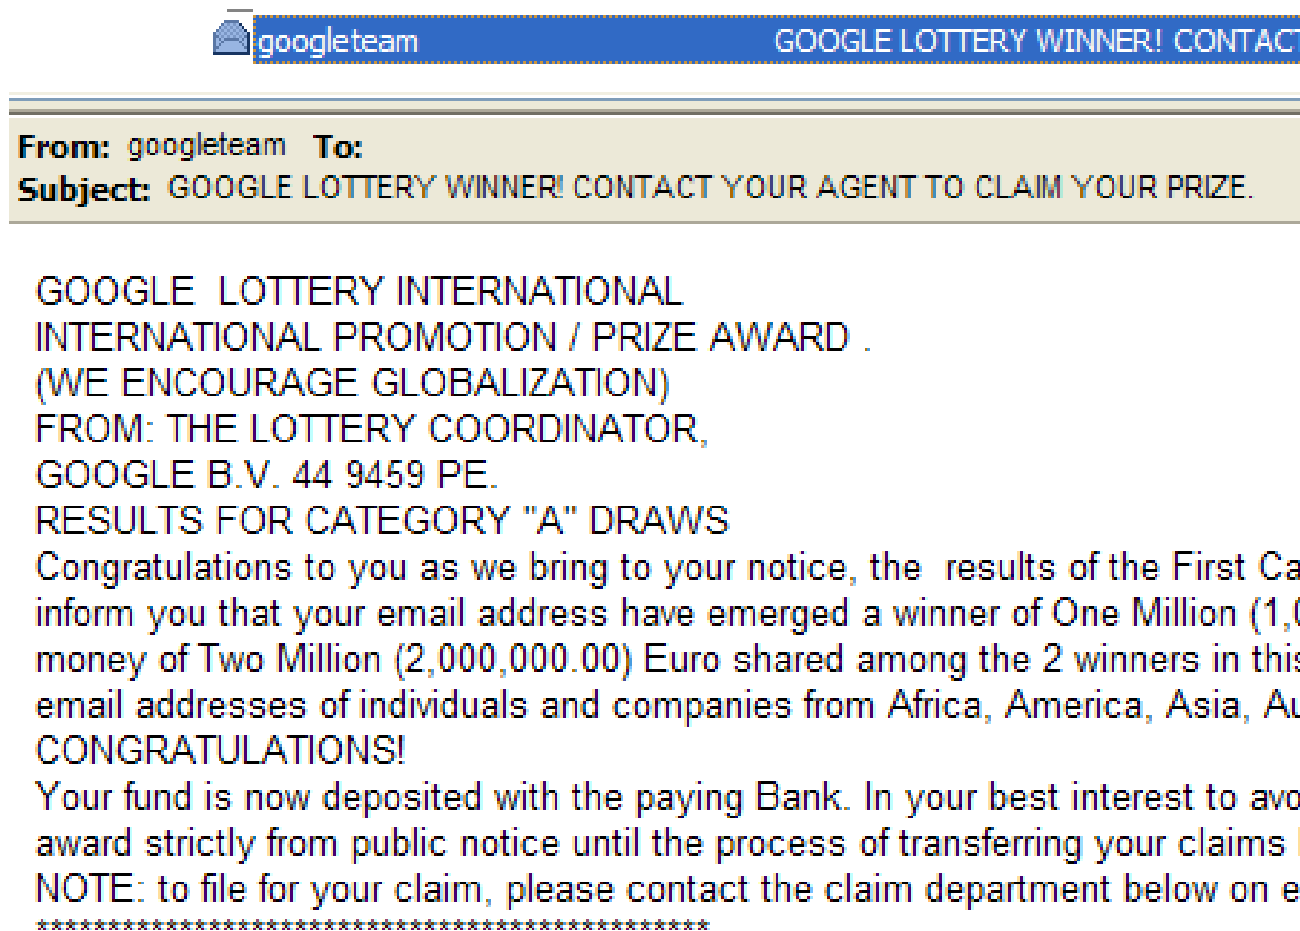
Machine Learning is Everywhere

- "If you invent a breakthrough in artificial intelligence, so machines can learn," Mr. Gates responded, "that is worth 10 Microsofts."

(Quoted in NY Times, Monday March 3, 2004)

Machine Learning is Everywhere

- Spam filtering



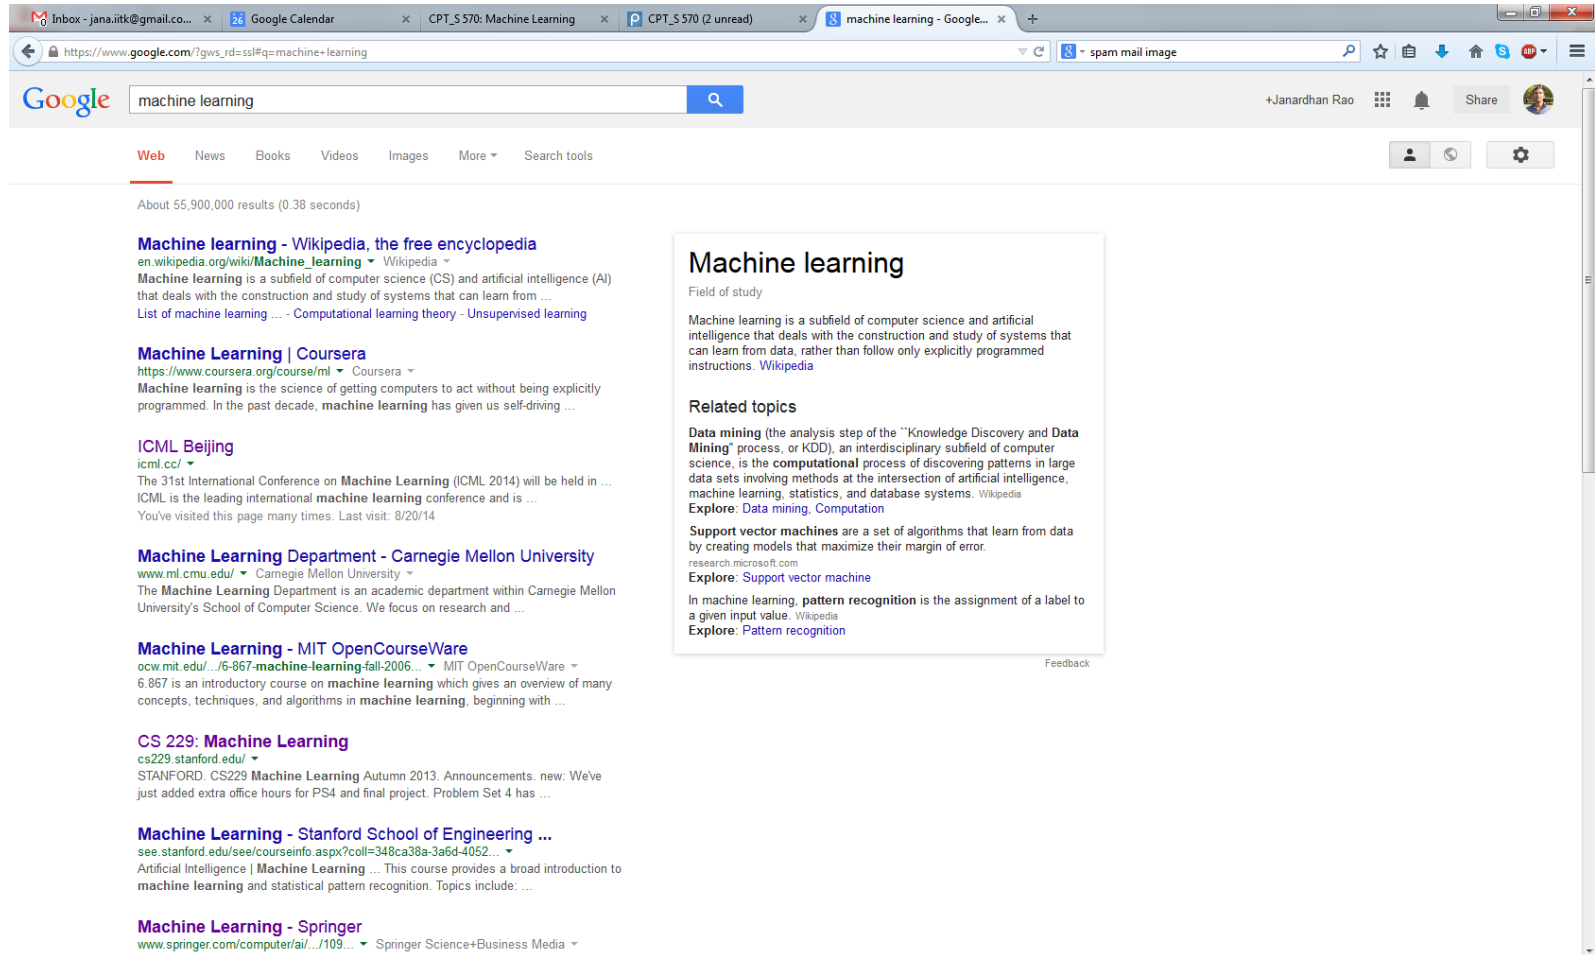
Machine Learning is Everywhere

- Optical Character Recognition (OCR)



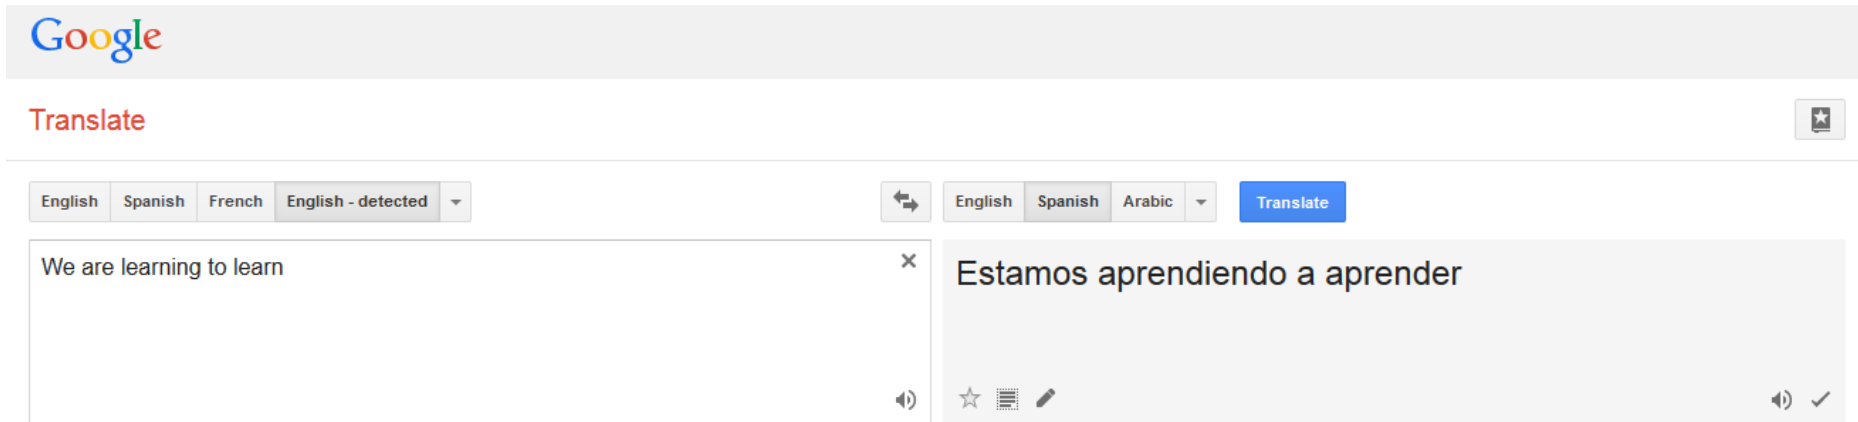
Machine Learning is Everywhere

- Search engines



Machine Learning is Everywhere

- Automatic Translation



Machine Learning is Everywhere

- Recommendation Engines

Related to Items You've Viewed

You viewed


Customers who viewed this also viewed




Machine Learning
► Tom M. Mitchell
Hardcover
★★★★☆ (48)
\$217.87



Learning From Data
► Hsuan-Tien Lin, Yaser S. Abu-Mostafa, Malik Magdon-Ismail
Hardcover
★★★★☆ (63)




Pattern Recognition and Machine Learning
► Christopher M. Bishop
Hardcover
★★★★☆ (97)
\$94.95 **\$71.44**



Machine Learning: The Art and Science...
► Peter A. Flach
Paperback
★★★★☆ (11)
\$64.00 **\$57.60**



The Elements of Statistical Learning...
Trevor Hastie, Robert Tibshirani, ...
Hardcover
★★★★☆ (36)
\$89.95 **\$71.96**



Machine Learning: An Algorithmic...
► Stephen Marsland
Hardcover
★★★★☆ (24)
\$79.95 **\$67.33**



Introduction to Machine Learning
► Ethem Alpaydin
Hardcover
★★★★☆ (26)
\$60.00 **\$48.12**

► [View or edit your browsing history](#)

Machine Learning is Everywhere

- Self-driving cars


Google's Self Driving Car for Blind People

by EDITORS on Apr 6, 2012 • 4:07 pm



ML Successes: Perception

Google Speech Recognition



Credit: Fernando Pereira & Matthew Firestone,
Google

Credit: Tom Dietterich

ML Successes: Image Captioning




"a black and white cat is sitting
on a chair."

Credit: Jeff Donahue, Trevor Darrell

ML Successes: Perception + Translation

Google Translate from Images




Credit: www.bbc.com

ML Successes: Skype Translator




credit: Skype

ML Successes: Reasoning (SAT)



Credit: Vijay Ganesh


ML Successes: Poker



Moore's Law

Credit: Michael Bowling

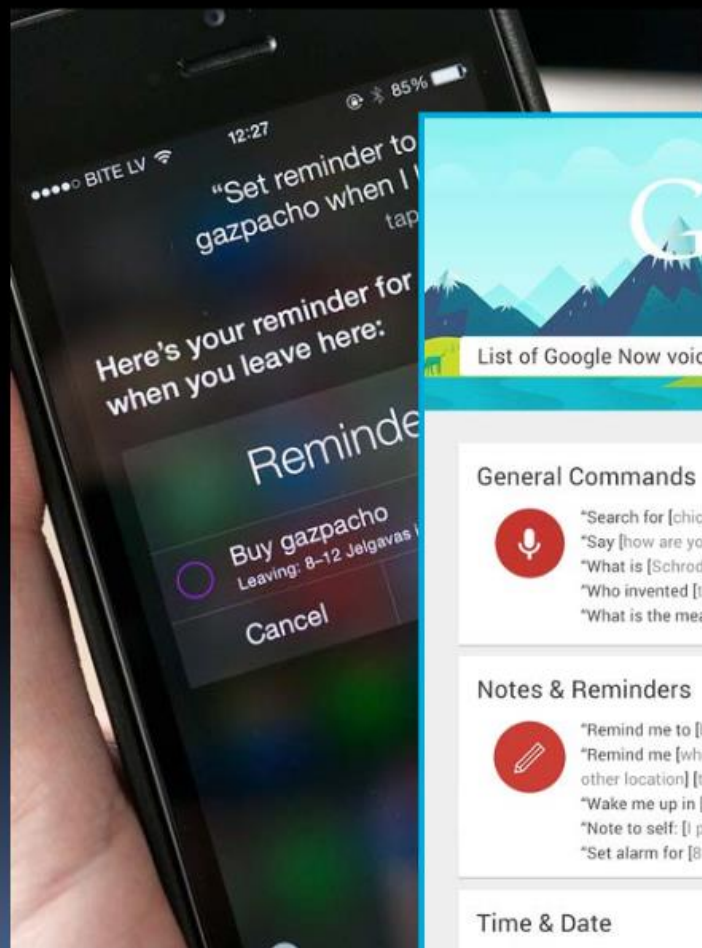
ML Successes: Chess and Go



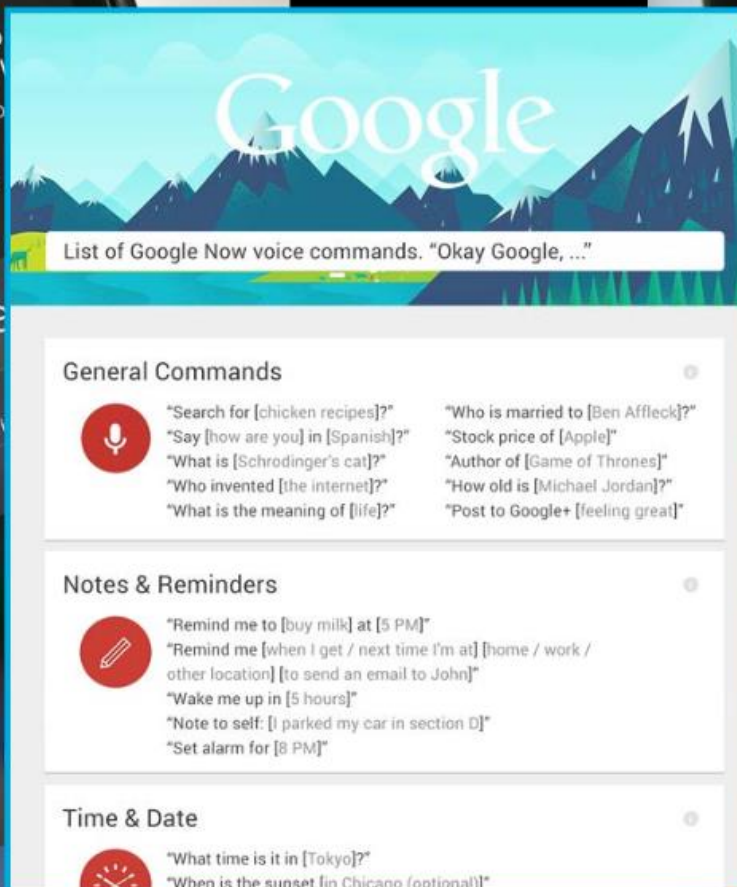
Silver, et al. (2016) *Nature*
Deep Learning +
Monte Carlo Tree Search

Credit: Martin Mueller


ML Successes: Personal Assistants



Credit: mashable.com



Credit: trendblog.net



Credit: The Verge

High-Stakes Applications: Self-Driving Cars



Credit: The Verge



Credit: delphi.com


Tesla AutoSteer



Credit: Tesla Motors

14

High-Stakes Applications: Automated Surgical Assistants



High-Stakes Applications: AI Hedge Funds



CADE METZ BUSINESS 01.25.16 7:00 AM

THE RISE OF THE ARTIFICIALLY INTELLIGENT HEDGE FUND

High-Stakes Applications: Power Grid Control

CONTROLLING THE POWER GRID WITH ARTIFICIAL INTELLIGENCE

02.07.2015

Credit: EBM Netz AG

DARPA Exploring Ways to Protect Nation's Electrical Grid from Cyber Attack

Effort calls for creation of automated systems to restore power within seven days or less after attack

Credit: DARPA

High-Stakes Applications: Autonomous Weapons

Northrop Grumman X-47B



Credit: Wikipedia

UK Brimstone Anti-Armor Weapon



Credit: Duch.seb - Own work, CC BY-SA 3.0

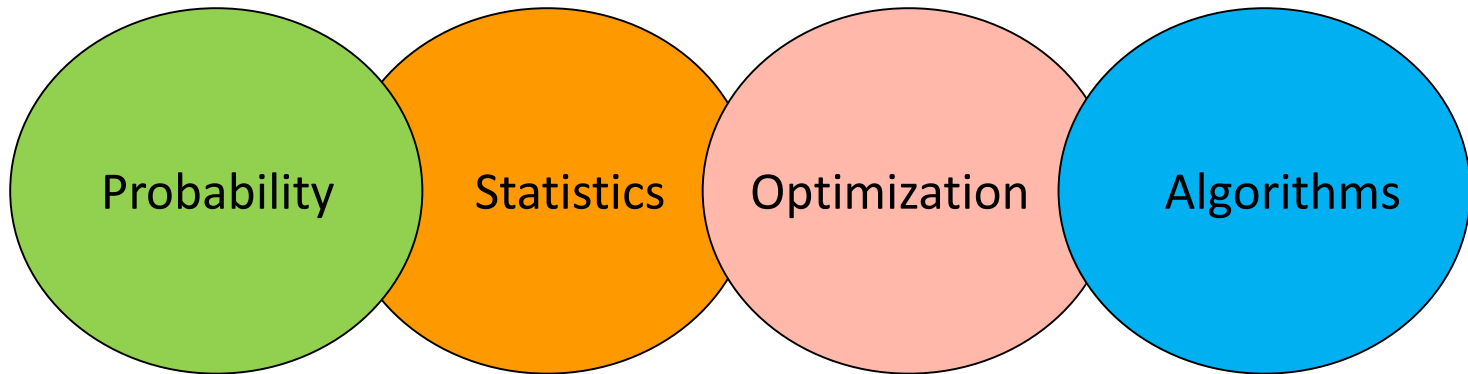
Samsung SGR-1



Credit: AFP/Getty Images

What is Machine Learning?


- Machine learning is the branch of engineering that develops technology for automated inference
 - ▲ It combines




What is Machine Learning?

- Machine learning = Automating Automation

Traditional Programming



Machine Learning




Learning Paradigms

- **Supervised Learning** – main focus of this course
- **Semi-Supervised Learning**
- **Active Learning**
- **Reinforcement Learning**

Supervised Learning

Learning a Classifier

(, male)


Training Data

$\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$

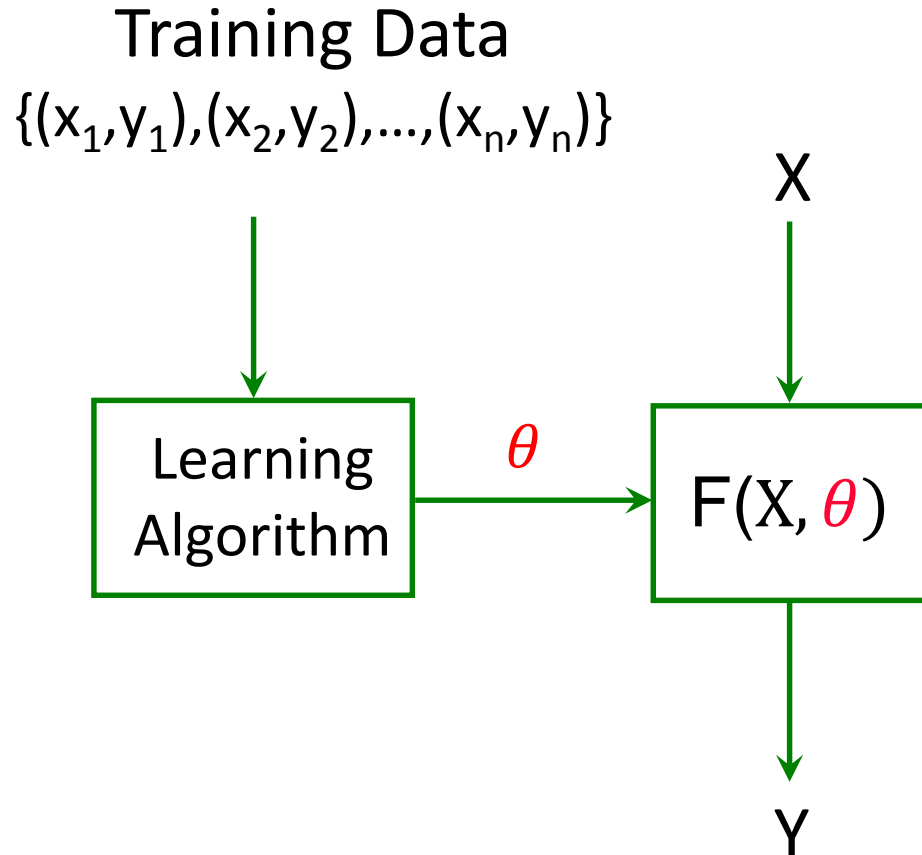
Example problem:

X - image of a face

$Y \in \{\text{male, female}\}$



Learning a Classifier

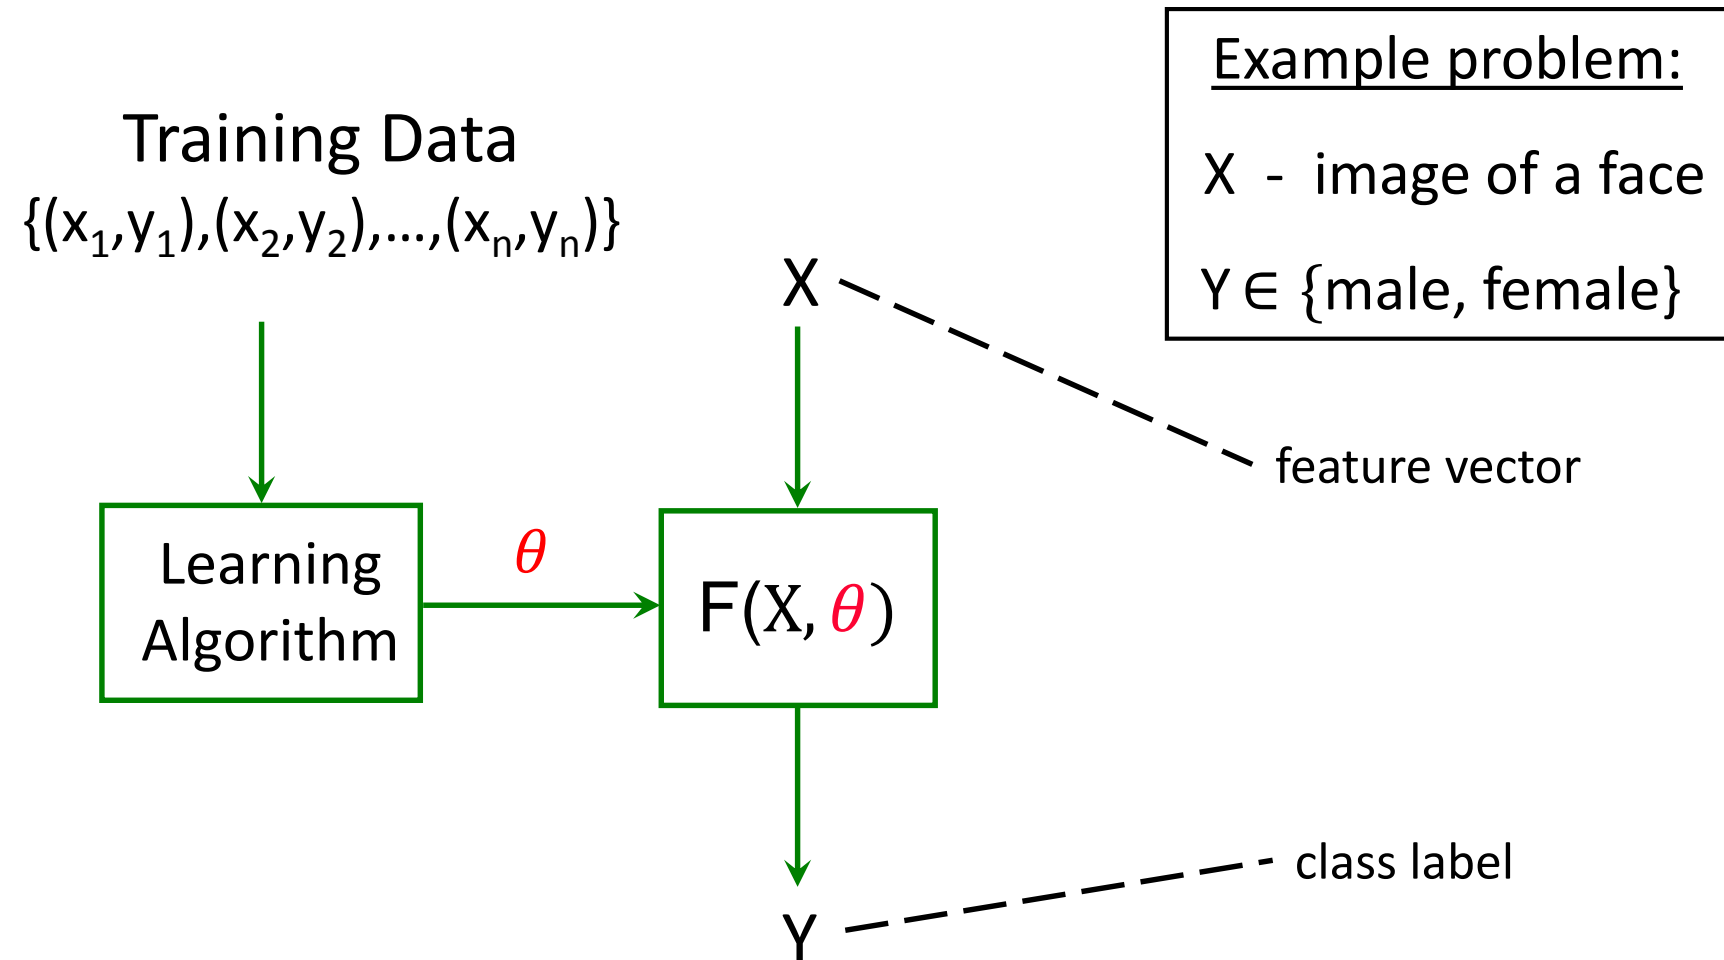


Example problem:

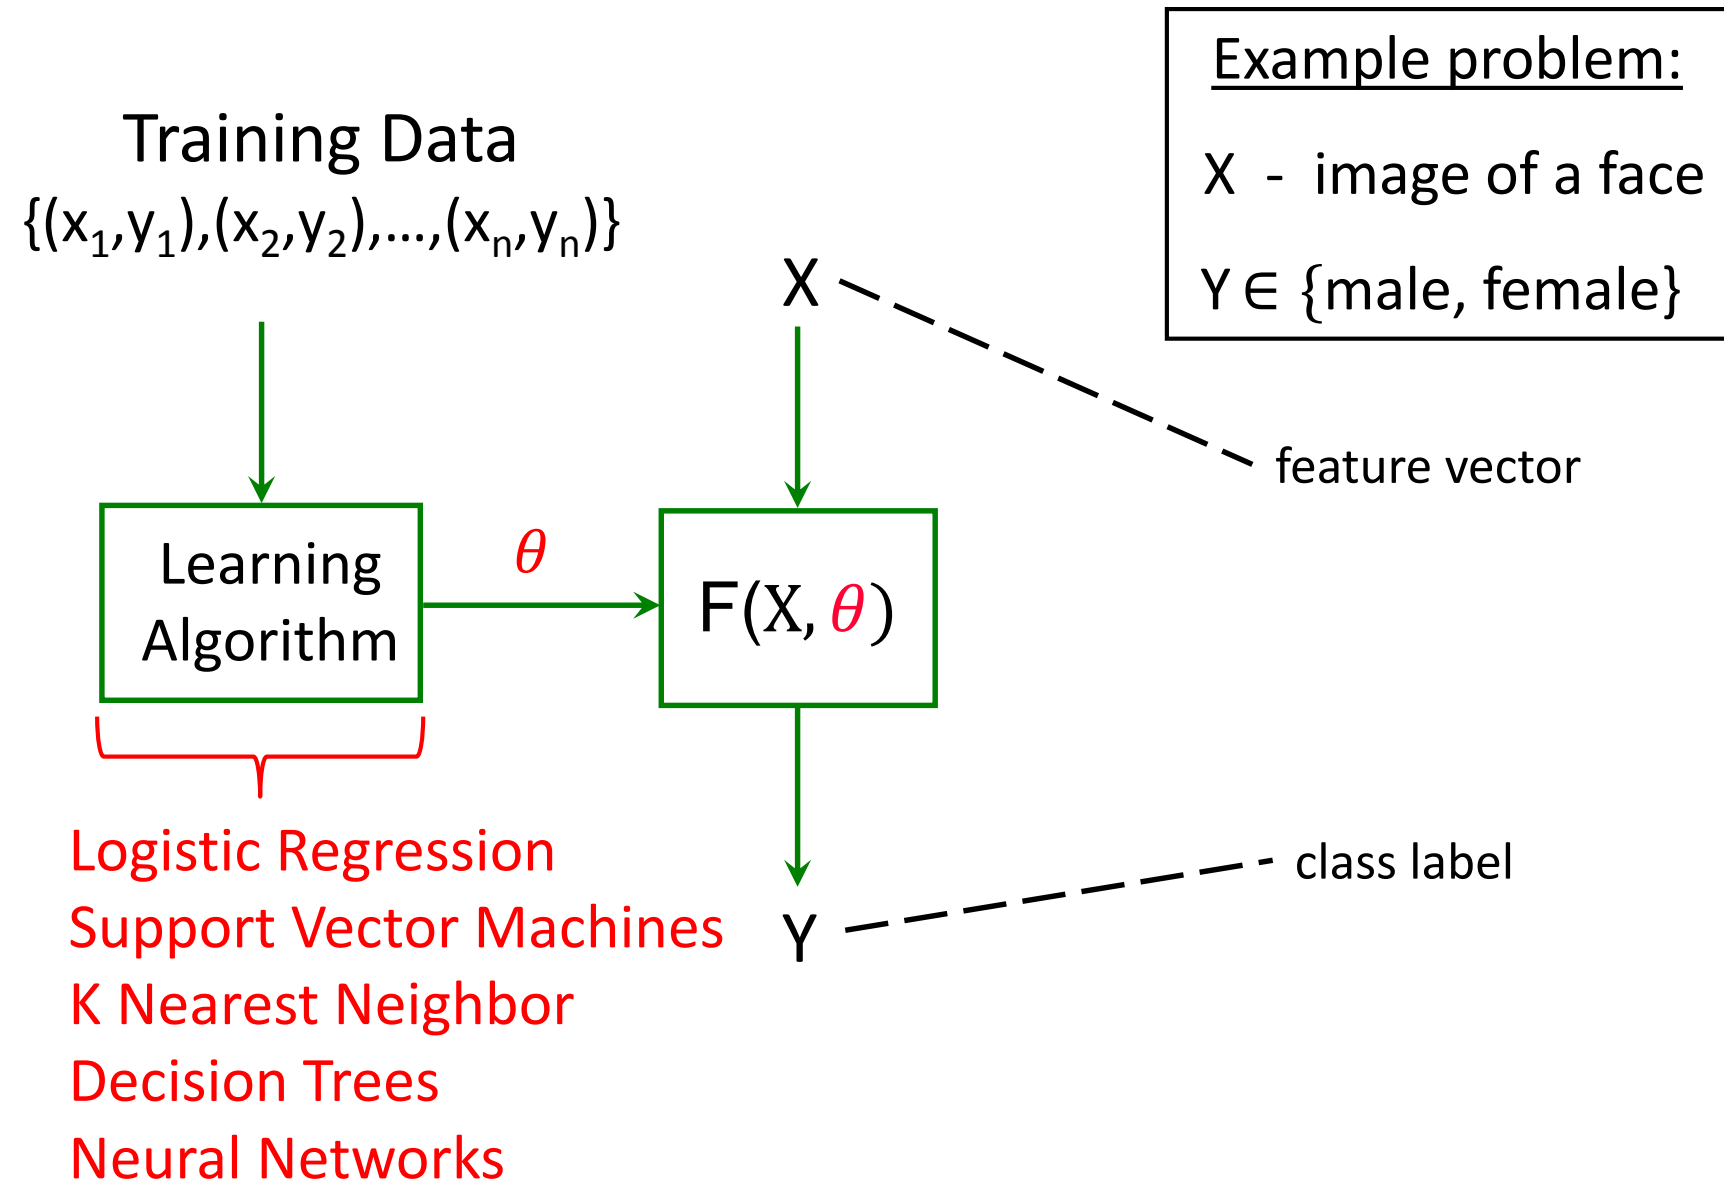
X - image of a face

$Y \in \{\text{male, female}\}$

Learning for Simple Outputs




Learning for Simple Outputs



Regression


- **Setting:** output y is a continuous value instead of a discrete value
 - ▲ Stock market price as a function of financial specs



Semi-Supervised Learning

Semi-Supervised Learning

- **Setting:** small amount of labeled data and large amount of unlabeled data



- ▲ find a classifier that separates the labeled points and separates the unlabeled points “well”

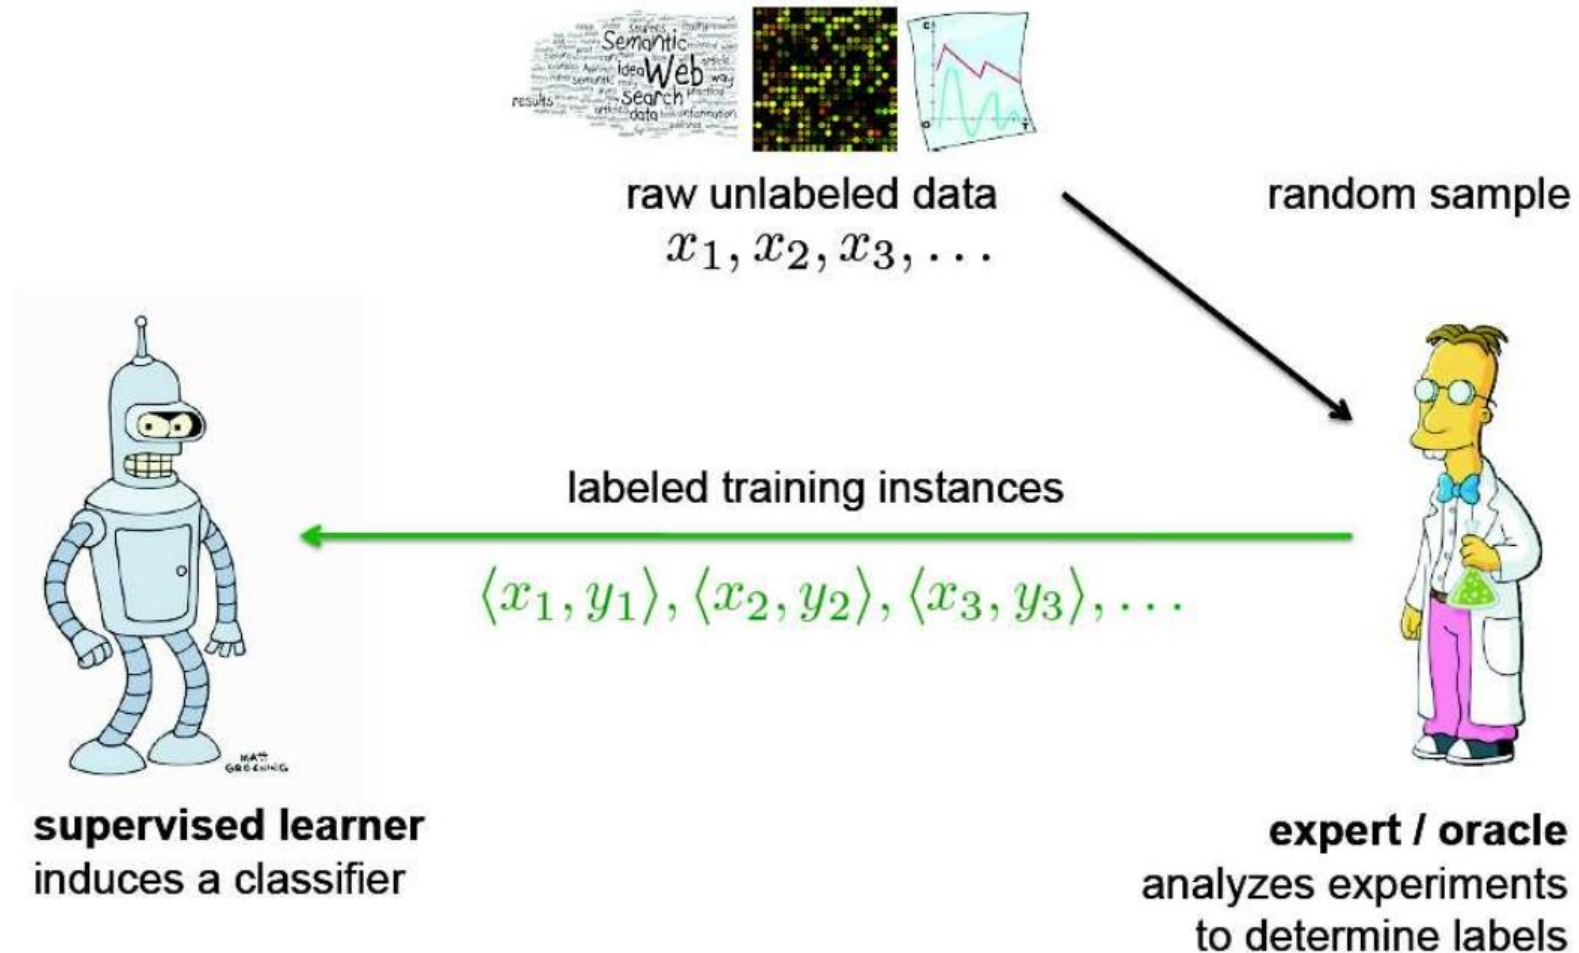
Semi-Supervised Learning

- **Co-Training Style Algorithms**

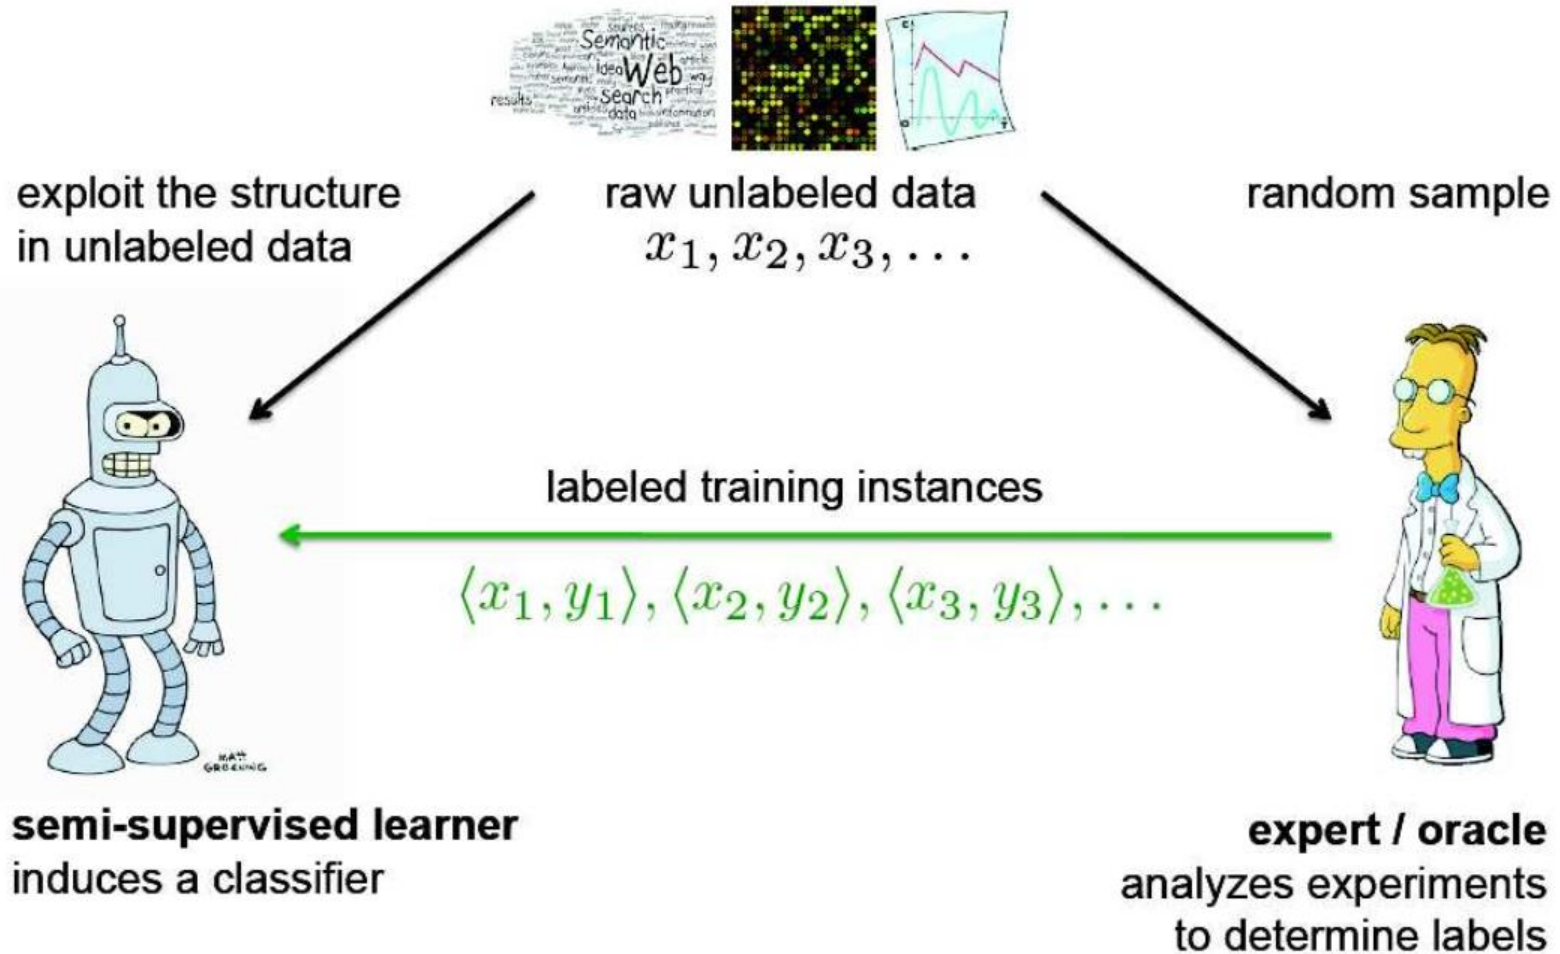
- ▲ Leverage diversity in the learners to learn from each other
- ▲ Diversity comes from multiple (redundant) views of the input – In webpage classification, one view is the “words” on the page and another view is the “links” that point to that page
- ▲ If only one view, employ learners with different hypothesis spaces to achieve diversity

Active Learning

(Passive) Supervised Learning



Semi-Supervised Learning




Active Learning



raw unlabeled data

x_1, x_2, x_3, \dots



Assumes some small
amount of initial labeled training data

active learner
induces a classifier




expert / oracle
analyzes experiments
to determine labels

Active Learning



inspect the
unlabeled data

raw unlabeled data
 x_1, x_2, x_3, \dots




active learner
induces a classifier



expert / oracle
analyzes experiments
to determine labels

Active Learning




inspect the
unlabeled data

raw unlabeled data
 x_1, x_2, x_3, \dots

request labels for selected data

$\langle x_1, ? \rangle$

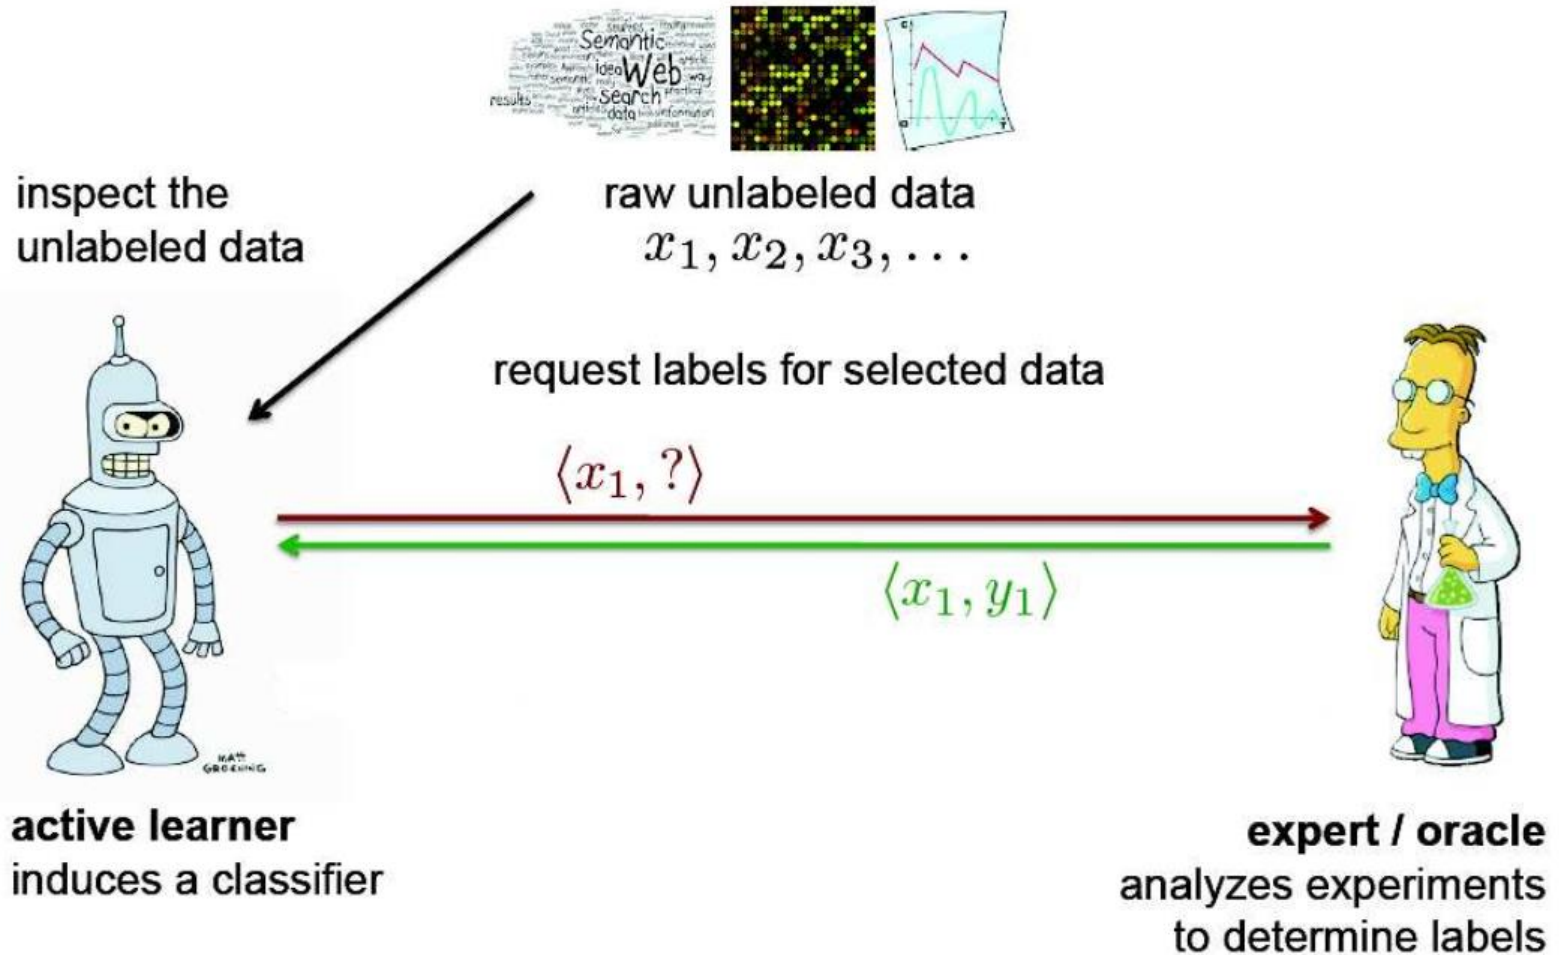


active learner
induces a classifier



expert / oracle
analyzes experiments
to determine labels

Active Learning



Active Learning



inspect the
unlabeled data


raw unlabeled data
 x_1, x_2, x_3, \dots

request labels for selected data

$\langle x_1, ? \rangle$

$\langle x_2, ? \rangle$

$\langle x_1, y_1 \rangle$



active learner
induces a classifier



expert / oracle
analyzes experiments
to determine labels

Active Learning



inspect the
unlabeled data

raw unlabeled data
 x_1, x_2, x_3, \dots


request labels for selected data

$\langle x_1, ? \rangle$

$\langle x_2, ? \rangle$

$\langle x_1, y_1 \rangle$

$\langle x_2, y_2 \rangle$



active learner
induces a classifier




expert / oracle
analyzes experiments
to determine labels

Motivation

- **Why do we need active learning?**
 - ▶ Supervised learning can solve all our problems, right?
 - ▶ Yes, if we have enough labeled data (input-output pairs)
 - ▶ But Labeling is expensive
 - ▶ We want to learn a highly-accurate function with few labeled examples
 - ▶ Intelligently select the examples for which we want to get labels for (unlabeled data is plentiful and cheap)

Active Learning Example: Drug Design

Goal: find compounds which bind to a particular target



Large collection of compounds, from:

- ▶ vendor catalogs
- ▶ corporate collections
- ▶ combinatorial chemistry

unlabeled point \equiv description of chemical compound
label \equiv *active* (binds to target) vs. *inactive*
getting a label \equiv chemistry experiment

Who uses Active Learning?



Sentiment analysis for blogs; Noisy relabeling
– *Prem Melville*



Biomedical NLP & IR; Computer-aided diagnosis
– *Balaji Krishnapuram*

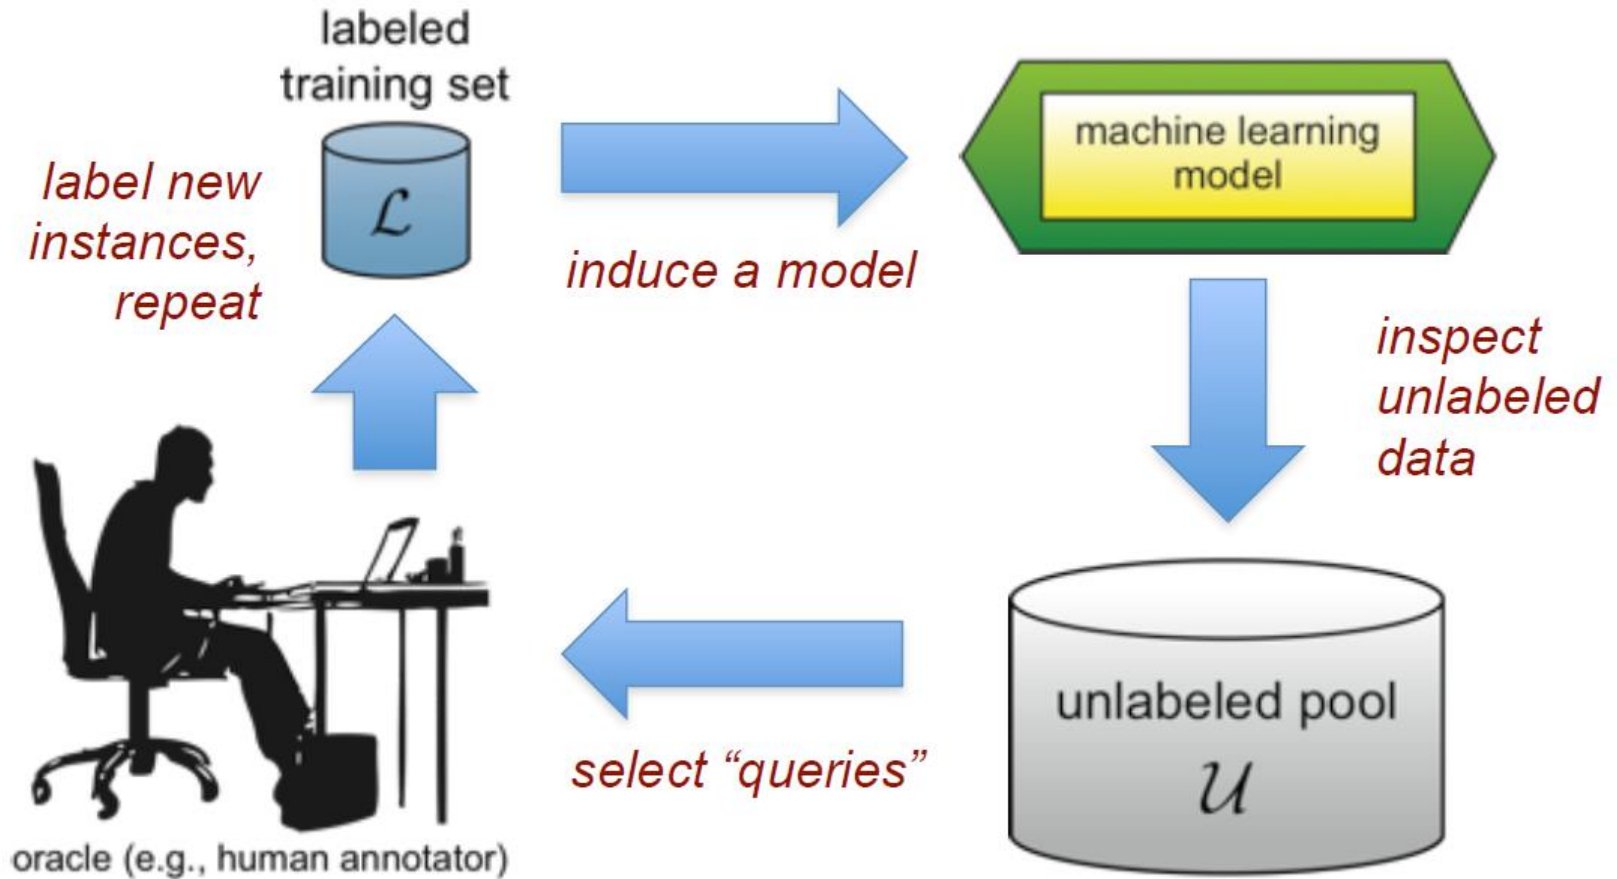


MS Outlook voicemail plug-in [Kapoor et al., IJCAI'07];
“A variety of prototypes that are in use throughout the company.” – *Eric Horvitz*



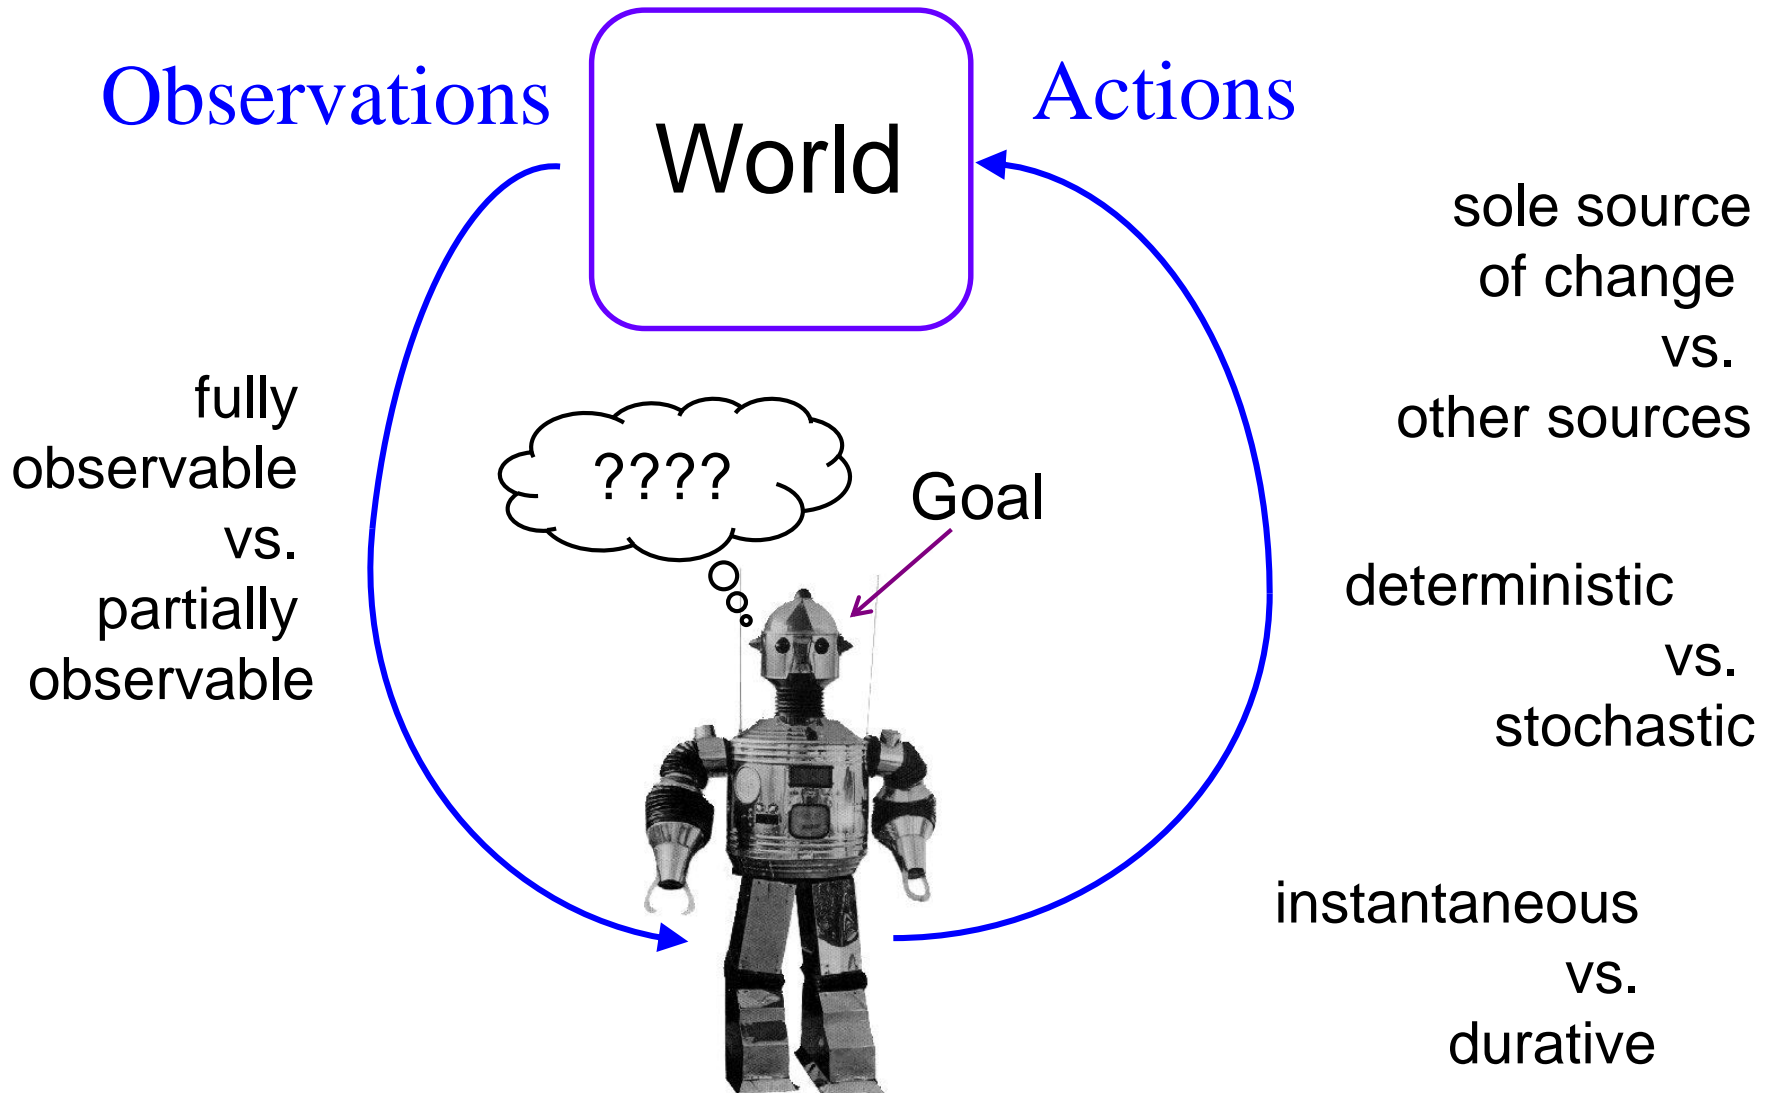
“While I can confirm that we're using active learning in earnest on many problem areas... I really can't provide any more details than that. Sorry to be so opaque!”
– *David Cohn*

Pool based Active Learning

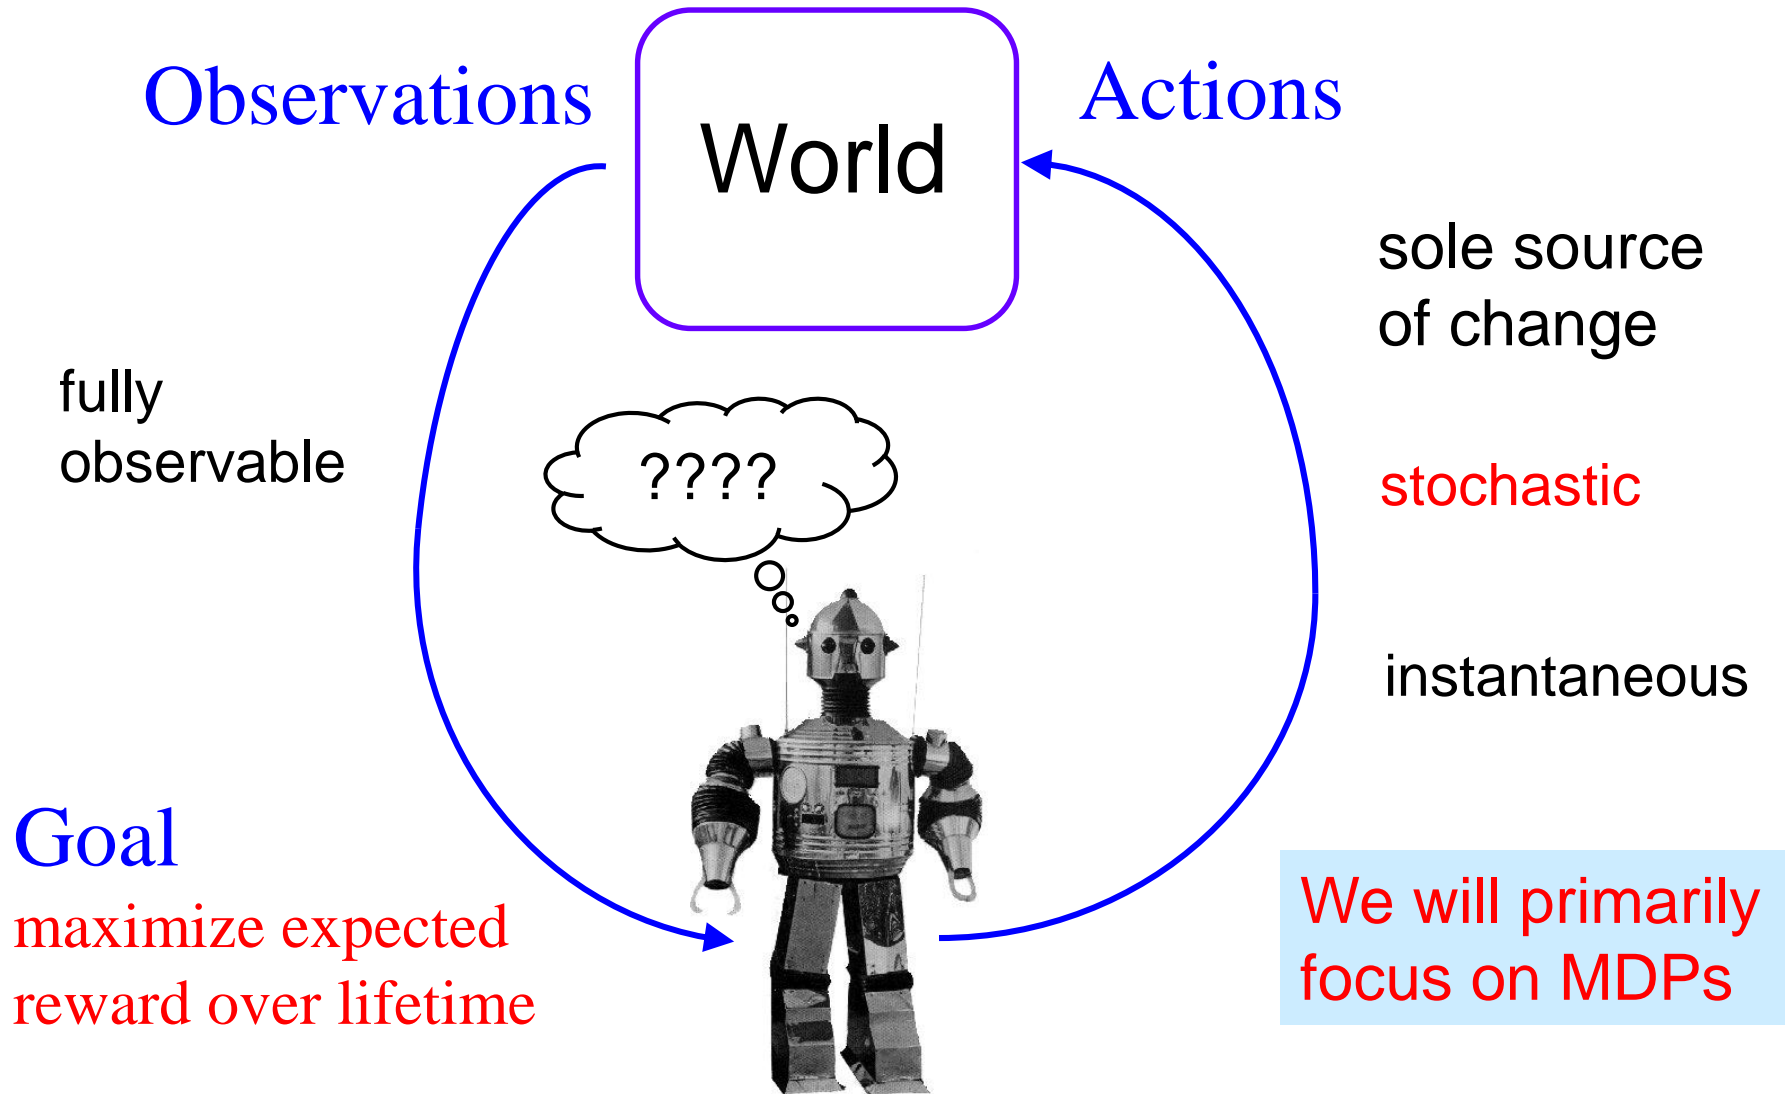


Reinforcement Learning


Reinforcement Learning



Stochastic/Probabilistic Planning: Markov Decision Process (MDP) Model



Stochastic/Probabilistic Planning: Markov Decision Process (MDP) Model



Example MDP



State describes
all visible info
about cards



Action are the
different legal
card movements

Goal


win the game or
play max # of cards




Input Representation and Abstract Machine Learning Algorithm

Learning for Simple Outputs

Training Data
 $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$



θ



Logistic Regression
Support Vector Machines
K Nearest Neighbor
Decision Trees
Neural Networks

X

??

feature vector

Y

class label

Example problem:

X - image of a face

$Y \in \{\text{male, female}\}$

Input examples: Representation

- In ML, our input examples (emails, text documents, images) are often represented as real-valued vectors:
 $x \in R^d$
 - ▲ each co-ordinate of x is called a **feature**
- Some examples
 - ▲ Bag-of-words representation of text
 - ▲ Histograms of colors in image
 - ▲ Sound frequency histogram


Input examples: Representation

- Bag-of-words model

- ▲ sentences to points


1. To be, or not to be,
2. To be a woman,
3. To not be a man

To	be	or	not	woman	a	man
2	2	1	1	0	0	0
1	1	0	0	1	1	0
1	1	0	1	0	1	1




Input examples: Representation

- Histogram of colors in image



Input examples: Representation

- Sound frequency histogram



Overview of ML Algorithms

- There are lot of machine learning algorithms
- Every machine learning algorithm has three components
 - ▲ **Representation**
 - ▲ **Evaluation**
 - ▲ **Optimization**

Representation: Examples

- Linear hyper-planes
- Decision trees
- Sets of conjunctive / logical rules
- Graphical models (Bayes/Markov nets)
- Neural Networks
- ...

Evaluation: Examples

- Accuracy
- Precision and recall
- Squared error
- Likelihood
- Cost / Utility
- Margin
- Entropy
- ...

Optimization: Examples

- **Combinatorial Optimization**
 - ▲ greedy search, dynamic programming
- **Convex Optimization**
 - ▲ gradient descent, co-ordinate descent
- **Constrained Optimization**
 - ▲ linear programming, quadratic programming
- ...

Machine Learned Programs: Errors

- **Approximation Error**
 - ▲ Error due to restricted hypothesis class (representation)
- **Estimation Error**
 - ▲ Error due to finite training samples
- **Optimization Error**
 - ▲ Error due to not finding a global optimum to the optimization problem

Learning Classifiers via Perceptron Algorithm

Formal setting – Classification

- **Instances**

- emails

$$\mathbf{x} \in \mathcal{X}$$

- **Labels**

- Spam vs. non-spam

$$y \in \mathcal{Y} = \{-1 ; 1\}$$

- **Prediction rule**

- Linear prediction rule

$$f(\mathbf{x}) = \hat{y}$$

- **Loss**

- No. of mistakes

$$\ell(\hat{y}, y) \in \mathbb{R}_+$$

Predictions

- Continuous predictions : $f : \mathcal{X} \rightarrow \mathbb{R}$
 - Label $\text{sign}(f(\mathbf{x}))$
 - Confidence $|f(\mathbf{x})|$
- Linear Classifiers
 - Prediction :
$$\begin{aligned}\hat{y} &= \text{sign}(f(\mathbf{x})) \\ &= \arg \max_{y \in \mathcal{Y}} \mathbf{w} \cdot \Phi(\mathbf{x}, y) \\ &= \text{sign}(\mathbf{w} \cdot \mathbf{x})\end{aligned}$$
$$|f(\mathbf{x})| = |\mathbf{w} \cdot \mathbf{x}|$$

Loss Functions

- **Natural Loss:**
 - Zero-One loss

$$\ell(\hat{y}, y) = \begin{cases} 0 & y = \hat{y} \\ 1 & y \neq \hat{y} \end{cases}$$

Online Framework

- Initialize Classifier $f_1(\mathbf{x})$
- Algorithm works in rounds $t = 1 \dots T \dots$
- On round t the online algorithm :
 - Receives an input instance \mathbf{x}_t
 - Outputs a prediction $f_t(\mathbf{x}_t) = \hat{y}_t$
 - Receives a feedback label y_t
 - Computes loss $\ell(\hat{y}_t, y_t)$
 - Updates the prediction rule $f_t \rightarrow f_{t+1}$
- Goal :
 - Suffer small cumulative loss $\sum_t \ell(\hat{y}_t, y_t)$

Why Online Learning?

- Fast
- Memory efficient - process one example at a time
- Simple to implement
- Formal guarantees – Mistake bounds
- Online to Batch conversions
- No statistical assumptions
- Adaptive

Update Rules

- Online algorithms are based on an update rule which defines f_{t+1} from f_t (and possibly other information)
- **Linear Classifiers** : find \mathbf{w}_{t+1} from \mathbf{w}_t based on the input (\mathbf{x}_t, y_t)
- **Perceptron algorithm employs a specific update rule**

Design Principle of Online Learning Algorithms

- If the learner suffers non-zero loss at any round, then we want to balance two goals:
 - **Corrective:** Change weights so that we **don't make this error again**
 - **Conservative:** **Don't change the weights too much**

The Perceptron Algorithm ($\eta = 1$)

- If No-Mistake $y_t(\mathbf{w}_t \cdot \mathbf{x}_t) > 0$

- Do nothing $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t$

- If Mistake $y_t(\mathbf{w}_t \cdot \mathbf{x}_t) \leq 0$

- Update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_t \mathbf{x}_t$

The Perceptron Algorithm ($\eta = 1$)

- When mistake happens, what does the update do?
 - **If $y_t = 1$:** $w_{t+1} = w_t + x_t$
 - ✓ w_{t+1} moves “closer to” x_t OR
 - ✓ x_t moves towards the positive side of the decision boundary
 - **If $y_t = -1$:** $w_{t+1} = w_t - x_t$
 - ✓ w_{t+1} moves “away from” x_t OR
 - ✓ x_t moves towards the negative side of the decision boundary
- In both cases, we are moving towards the “correct solution”

The Perceptron Algorithm

- Suppose w_t makes a mistake on (x_t, y_t) , and we update w_{t+1} as $w_{t+1} = w_t + y_t x_t$. Is it possible for w_{t+1} to also make a mistake on (x_t, y_t) ?



The Perceptron Algorithm ($\eta = 1$)

- Suppose w_t makes a mistake on (x_t, y_t) , and we update w_{t+1} as $w_{t+1} = w_t + y_t x_t$. Is it possible for w_{t+1} to also make a mistake on (x_t, y_t) ?




- ▲ Yes, depends on the learning rate η


When does Perceptron converge?

- **Linear Separability**

- ▲ There exists a hyper-plane (weight vector) separating the positive and negative points



Linearly separable



Not linearly separable

Measure of Separability


- **Margin**

- ▲ For a weight vector \mathbf{w} , and training set \mathcal{S} , margin of \mathbf{w} with respect to \mathcal{S} is defined as follows:


$$\gamma(\mathbf{w}) = \min_{(x,y) \in \mathcal{S}} y(\mathbf{w} \cdot \mathbf{x})$$

- The training data \mathcal{S} is linearly separable if there exists *at least* one weight vector \mathbf{w} for which the margin is positive, i.e., $\gamma(\mathbf{w}) > 0$.

Margin: Examples



Low margin data



High margin data

Perceptron: Convergence Result

- **Theorem:** If the training data is linearly separable with margin γ , and if $\|x_i\| \leq 1$ for all examples (x_i, y_i) in the training set, then perceptron makes $\leq \frac{1}{\gamma^2}$ mistakes.
 - ▲ Proof??
- Lower margin implies more mistakes
- May need more than one pass over the training data to get a classifier with no mistakes

What if data is not linearly separable?

- Ideally, we want to find a linear separator that makes the minimum number of mistakes on the training data
 - ▲ NP-Hard problem! (Minsky and Papert, 1969)
 - ▲ This result killed the neural networks research in 1970's
- Perceptron still works
 - ▲ there will be few mistakes close to the decision boundary
 - ▲ will never converge to a single w as we make more passes

Problems with Perceptron

- Doesn't converge with inseparable data
 - Weight updates may often be very “bold”
 - Doesn't optimize margin
 - Sensitive to the order of examples
- ▲ **Voted and Averaged perceptron**

Voted Perceptron

- **Initialization:** $m = 1$; $w_1 = 0$; $c_m = 1$
- **Training Examples:** for $t = 1, 2, 3, \dots$
 - ▲ If mistake, update weights
 - $w_{m+1} = w_m + y_t x_t$
 - $m = m + 1$
 - $c_m = 1$
 - ▲ Else
 - $c_m = c_m + 1$ // counting how long w_m survived
- **Output:** $(w_1, c_1), (w_2, c_2), (w_3, c_3), \dots$

Voted Perceptron Classifier

$$f(x) = \textit{sign} \left(\sum_{i=1}^m c_i \textit{sign}(< w_i, x >) \right)$$

- Any drawbacks of voted perceptron?

Voted Perceptron Classifier

$$f(x) = \text{sign} \left(\sum_{i=1}^m c_i \text{sign}(< w_i, x >) \right)$$

- Any drawbacks of voted perceptron?
- Yes, we have to store all the classifiers (in practice could be many)
- How can we solve this problem?

Averaged Perceptron

- Same algorithm as voted perceptron, but the classification rule is different

$$f_{average}(x) = \text{sign} \left(\sum_{i=1}^m (< c_i w_i, x >) \right)$$

$$f_{voted}(x) = \text{sign} \left(\sum_{i=1}^m c_i \text{sign}(< w_i, x >) \right)$$

Averaged vs. Voted Perceptron

- Simple Example: If $c_1 = c_2 = c_3 = 1$

$$f_{average}(x) = \textit{sign}(\langle w_1 + w_2 + w_3, x \rangle)$$

$$f_{voted}(x) = \textit{majority sign of } \langle w_1, x \rangle, \langle w_2, x \rangle, \langle w_3, x \rangle$$

Some Practical Tricks

- **Shuffling**

- ▶ shuffling the training examples in each iteration

- **Variable learning rate**

- ▶ decrease as learning progresses
- ▶ follow some schedule
- ▶ Set automatically by line search (converges faster)
- ▶ See Leon Bottou's SGD website:
<http://leon.bottou.org/projects/sgd>

- **Averaged Perceptron can be implemented very efficiently** (See Algorithm 7 in Hal's chapter)

Some Practical Tricks

- **Learning Curve**

- ▶ Training iterations vs. number of mistakes
- ▶ You want to see that the mistakes decrease as we increase the no. of iterations (curve goes down)
- ▶ Very useful in debugging and seeing the behavior of online learning algorithms

- **Hyper-parameter Optimization**

- ▶ Split the training data: sub-train + validation data
- ▶ Tune hyper-parameters (e.g., no. of iterations) on the validation data
- ▶ The learner should not look at the test data!

Multi-Class Classification: Setup

Suppose we have $(K > 2)$ classes.

K weight vectors: $w_1, w_2, \dots, w_K \in \mathbb{R}^c$

input instance $x \in \mathbb{R}^d$

$$\text{Score}(\text{label } r) = w_r \cdot x$$

Multi-Class Classification: Learning

Class y	$w_y \cdot x$
1	-1.08
2	1.66
3	0.37
4	-2.09

Prediction: output label (class) with highest score.

Learning:

$$w_{y^*} = w_{y^*} + x$$
$$w_{\hat{y}} = w_{\hat{y}} - x$$

Regression Learning: Setup

Regression Learning:

y is continuous value.

Prediction Rule: $F(x) = w \cdot x$

Widrow-Hoff Algorithm:

- Initialize $w_1 = 0$

for $t = 1$ to T do

- get $x_t \in \mathbb{R}^d$

- predict $\hat{y}_t = w_t \cdot x_t$

- observe y_t^*

- Incur loss of $(\hat{y}_t - y_t^*)^2$

- $w_{t+1} = w_t - \eta (w_t \cdot x_t - y_t^*) x_t$

end